

# Comprehensive Generative AI Learning Path

Curated by Anish Roychowdhury

April 4, 2025

## Contents

<b>1</b>	<b>Foundations of Generative AI</b>	<b>3</b>
1.1	Introduction to Generative AI . . . . .	3
1.2	Large Language Models (LLMs) Basics . . . . .	3
<b>2</b>	<b>Prompt Engineering</b>	<b>3</b>
2.1	Basics of Prompt Engineering . . . . .	3
2.2	Advanced Prompt Engineering . . . . .	3
<b>3</b>	<b>Generative Models</b>	<b>3</b>
3.1	Diffusion Models . . . . .	3
3.2	GANs (Generative Adversarial Networks) . . . . .	4
<b>4</b>	<b>Building AI Agents</b>	<b>4</b>
4.1	AI Agent Frameworks . . . . .	4
4.2	Multi-Agent Systems . . . . .	4
<b>5</b>	<b>Retrieval-Augmented Generation (RAG)</b>	<b>4</b>
5.1	Basics of RAG . . . . .	4
5.2	Vector Databases . . . . .	4
<b>6</b>	<b>Fine-Tuning LLMs</b>	<b>5</b>
6.1	Parameter-Efficient Fine-Tuning (PEFT) . . . . .	5
6.2	Instruction Fine-Tuning . . . . .	5
<b>7</b>	<b>Reinforcement Learning from Human Feedback (RLHF)</b>	<b>5</b>
7.1	Introduction to RLHF . . . . .	5
<b>8</b>	<b>Model Context Protocol (MCP)</b>	<b>5</b>
8.1	Fundamentals of MCP . . . . .	5
8.2	Implementing MCP in Applications . . . . .	5
8.3	Advanced MCP Patterns . . . . .	6
<b>9</b>	<b>Deployment of Generative AI Models</b>	<b>6</b>
9.1	Local Deployment . . . . .	6
9.2	Serverless Deployment . . . . .	6
9.3	Edge Deployment . . . . .	6
<b>10</b>	<b>Evaluation Metrics for Generative AI Models</b>	<b>6</b>
10.1	Error Metrics and Loss Computation . . . . .	6
10.2	Advanced Evaluation Frameworks . . . . .	6

<b>11 Practical Projects and Applications</b>	<b>7</b>
11.1 End-to-End Gen AI Projects . . . . .	7
11.2 Industry-Specific Applications . . . . .	7
<b>12 Ethical Considerations and Responsible AI</b>	<b>7</b>
12.1 Ethical Frameworks . . . . .	7
12.2 Bias Mitigation . . . . .	7
<b>13 Future Directions in Generative AI</b>	<b>7</b>
13.1 Multimodal Models . . . . .	7
13.2 AI Research Frontiers . . . . .	8
<b>14 Community Resources and Continuous Learning</b>	<b>8</b>
14.1 Open Source Communities . . . . .	8
14.2 Conferences and Workshops . . . . .	8

# 1 Foundations of Generative AI

## 1.1 Introduction to Generative AI

- Overview of generative models and applications.
  - [DeepLearning.AI Generative AI for Beginners](https://github.com/microsoft/generative-ai-for-beginners)  
<https://github.com/microsoft/generative-ai-for-beginners>
  - [Big Book of Generative AI by Databricks](https://www.databricks.com/resources/ebook/big-book-generative-ai)  
<https://www.databricks.com/resources/ebook/big-book-generative-ai>

## 1.2 Large Language Models (LLMs) Basics

- Learn about LLMs, tokenization, embeddings, and transformers.
  - [Building LLMs from Scratch \(15 Lectures\)](https://lnkd.in/gifp9PkM)  
<https://lnkd.in/gifp9PkM>
  - [Transformer Architecture Explained](https://jalammar.github.io/illustrated-transformer/)  
<https://jalammar.github.io/illustrated-transformer/>
  - [Transformers from Scratch in Python](https://towardsdatascience.com/build-your-own-transformer-from-scratch-using-pytorch-84c850470dc6)  
<https://towardsdatascience.com/build-your-own-transformer-from-scratch-using-pytorch-84c850470dc6>

# 2 Prompt Engineering

## 2.1 Basics of Prompt Engineering

- Techniques like zero-shot, one-shot, few-shot learning, and chain-of-thought prompting.
  - [Prompt Engineering Guide](https://aman.ai/primers/ai/prompt-engineering/)  
<https://aman.ai/primers/ai/prompt-engineering/>
  - [LangChain Prompt Templates](https://lnkd.in/dVkuizQ)  
<https://lnkd.in/dVkuizQ>

## 2.2 Advanced Prompt Engineering

- Generated knowledge, prompt chaining, and ReAct.
  - [Prompt Tuning with Hugging Face](https://huggingface.co/docs/peft/task_guides/clm-prompt-tuning)  
[https://huggingface.co/docs/peft/task\\_guides/clm-prompt-tuning](https://huggingface.co/docs/peft/task_guides/clm-prompt-tuning)
  - [NVIDIA Blog on Prompt Engineering](https://developer.nvidia.com/blog/an-introduction-to-large-language-models-prompt-engineering-and-reasoning/)  
<https://developer.nvidia.com/blog/an-introduction-to-large-language-models-prompt-engineering-and-reasoning/>

# 3 Generative Models

## 3.1 Diffusion Models

- Learn how diffusion models generate images and other content.
  - [Diffusion Models Explained](https://towardsdatascience.com/diffusion-models-made-easy-8414298ce4da)  
<https://towardsdatascience.com/diffusion-models-made-easy-8414298ce4da>
  - [Illustrated Stable Diffusion](https://jalammar.github.io/illustrated-stable-diffusion/)  
<https://jalammar.github.io/illustrated-stable-diffusion/>

## 3.2 GANs (Generative Adversarial Networks)

- Explore GAN architectures and applications.
  - [Understanding GANs](https://towardsdatascience.com/understanding-generative-adversarial-networks-gans-cd6e4651a29)  
<https://towardsdatascience.com/understanding-generative-adversarial-networks-gans-cd6e4651a29>
  - [GAN Lecture Series](https://lnkd.in/eSf66zT)  
<https://lnkd.in/eSf66zT>

## 4 Building AI Agents

### 4.1 AI Agent Frameworks

- Learn frameworks like LangChain, CrewAI, and AutoGen.
  - [LangChain Crash Course](https://medium.com/databutton/getting-started-with-langchain-a-powerful-tool-for-working-with-lar)  
<https://medium.com/databutton/getting-started-with-langchain-a-powerful-tool-for-working-with-lar>
  - [CrewAI GitHub Repository](https://github.com/joaomdmoura/crewAI)  
<https://github.com/joaomdmoura/crewAI>

### 4.2 Multi-Agent Systems

- Practical workflows for multi-agent environments.
  - [Multi AI Agent Systems with CrewAI](https://lnkd.in/dTudrD55)  
<https://lnkd.in/dTudrD55>
  - [OpenAI Agents SDK in Python](https://github.com/openai/openai-agents-python)  
<https://github.com/openai/openai-agents-python>

## 5 Retrieval-Augmented Generation (RAG)

### 5.1 Basics of RAG

- Combine LLMs with external knowledge bases using vector databases.
  - [Guide to RAG Frameworks](https://deepchecks.com/practical-guide-to-crafting-your-first-llm-powered-app-using-rag-framework)  
<https://deepchecks.com/practical-guide-to-crafting-your-first-llm-powered-app-using-rag-framework>
  - [LangChain for RAG Applications](https://medium.com/mlearning-ai/create-a-chatbot-in-python-with-langchain-and-rag-85bfba8c62d2)  
<https://medium.com/mlearning-ai/create-a-chatbot-in-python-with-langchain-and-rag-85bfba8c62d2>

### 5.2 Vector Databases

- Learn about vector search and embedding storage solutions.
  - [Introduction to Vector Databases](https://weaviate.io/blog/what-is-a-vector-database)  
<https://weaviate.io/blog/what-is-a-vector-database>
  - [FAISS Tutorial](https://medium.com/mlearning-ai/mastering-similarity-searches-building-a-faiss-index-with-cosine-)  
<https://medium.com/mlearning-ai/mastering-similarity-searches-building-a-faiss-index-with-cosine->

## 6 Fine-Tuning LLMs

### 6.1 Parameter-Efficient Fine-Tuning (PEFT)

- Techniques like LoRA and QLoRA for efficient tuning.
  - [Fine-Tuning with LoRA](https://abvijaykumar.medium.com/fine-tuning-llm-parameter-efficient-fine-tuning-peft-lora-qlora-p)  
<https://abvijaykumar.medium.com/fine-tuning-llm-parameter-efficient-fine-tuning-peft-lora-qlora-p>
  - [Hugging Face Blog on PEFT](https://huggingface.co/blog/trl-peft)  
<https://huggingface.co/blog/trl-peft>

### 6.2 Instruction Fine-Tuning

- Customize LLMs for specific tasks using instruction tuning.
  - [Instruction Fine-Tuning Guide](https://medium.com/@ud.chandra/instruction-fine-tuning-llama-2-with-pefts-qlora-method-d6a801ebb1)  
<https://medium.com/@ud.chandra/instruction-fine-tuning-llama-2-with-pefts-qlora-method-d6a801ebb1>
  - [Fine-Tune Your Own Llama Model in Colab](https://towardsdatascience.com/fine-tune-your-own-llama-2-model-in-a-colab-notebook-df9823a04a32)  
<https://towardsdatascience.com/fine-tune-your-own-llama-2-model-in-a-colab-notebook-df9823a04a32>

## 7 Reinforcement Learning from Human Feedback (RLHF)

### 7.1 Introduction to RLHF

- Understand how RLHF improves model alignment with human preferences.
  - [Hugging Face RLHF Blog](https://huggingface.co/blog/rlhf)  
<https://huggingface.co/blog/rlhf>
  - [Short Course on RLHF by DeepLearning.AI](https://www.deeplearning.ai/short-courses/reinforcement-learning-from-human-feedback/)  
<https://www.deeplearning.ai/short-courses/reinforcement-learning-from-human-feedback/>

## 8 Model Context Protocol (MCP)

### 8.1 Fundamentals of MCP

- Understanding the Model Context Protocol for LLM interactions.
  - [Anthropic's MCP Guide](https://github.com/anthropics/anthropic-cookbook/blob/main/model_context_protocol/model_context_protocol.md)  
[https://github.com/anthropics/anthropic-cookbook/blob/main/model\\_context\\_protocol/model\\_context\\_protocol.md](https://github.com/anthropics/anthropic-cookbook/blob/main/model_context_protocol/model_context_protocol.md)
  - [Claude MCP Documentation](https://docs.anthropic.com/claude/docs/model-context-protocol)  
<https://docs.anthropic.com/claude/docs/model-context-protocol>

### 8.2 Implementing MCP in Applications

- Practical implementations and best practices.
  - [Structured Outputs with Claude 3 MCP](https://medium.com/@saitejasagi/structured-outputs-with-anthropics-claude-3-model-context-protocol)  
<https://medium.com/@saitejasagi/structured-outputs-with-anthropics-claude-3-model-context-protocol>
  - [OpenAI's Approach to Structured Outputs](https://github.com/openai/openai-cookbook/blob/main/examples/How_to_format_inputs_to_GPT_3_5_turbo_models.md)  
[https://github.com/openai/openai-cookbook/blob/main/examples/How\\_to\\_format\\_inputs\\_to\\_GPT\\_3\\_5\\_turbo\\_models.md](https://github.com/openai/openai-cookbook/blob/main/examples/How_to_format_inputs_to_GPT_3_5_turbo_models.md)

## 8.3 Advanced MCP Patterns

- Complex data extraction and validation patterns.
  - [Anthropic Research on MCP](https://www.anthropic.com/research/claude-3-model-context-protocol)  
<https://www.anthropic.com/research/claude-3-model-context-protocol>
  - [MCP Implementation Examples](https://github.com/ArnaudBuchholz/reserve-mcp)  
<https://github.com/ArnaudBuchholz/reserve-mcp>

# 9 Deployment of Generative AI Models

## 9.1 Local Deployment

- Run models locally on your machine.
  - [5 Free Tools for Local Deployment of LLMs](https://lnkd.in/dJsRrn2c)  
<https://lnkd.in/dJsRrn2c>

## 9.2 Serverless Deployment

- Deploy serverless applications using platforms like AWS Bedrock.
  - [Serverless Workflows with Amazon Bedrock](https://lnkd.in/dENCd795)  
<https://lnkd.in/dENCd795>

## 9.3 Edge Deployment

- Bring generative AI capabilities to edge devices.
  - [Deploying LLMs at the Edge with NVIDIA IGX Orin Developer Kit](https://lnkd.in/d94BxVjw)  
<https://lnkd.in/d94BxVjw>

# 10 Evaluation Metrics for Generative AI Models

## 10.1 Error Metrics and Loss Computation

- Evaluate model performance using metrics like BLEU score and perplexity.
  - [Foundations of NLP Metrics: BLEU Score and WER Metrics](https://towardsdatascience.com/foundations-of-nlp-explained-bleu-score-and-wer-metrics-1a5ba06d81)  
<https://towardsdatascience.com/foundations-of-nlp-explained-bleu-score-and-wer-metrics-1a5ba06d81>
  - [Perplexity in Language Models Explained](https://medium.com/@priyankads/perplexity-of-language-models-41160427ed72)  
<https://medium.com/@priyankads/perplexity-of-language-models-41160427ed72>

## 10.2 Advanced Evaluation Frameworks

- Comprehensive evaluation strategies for generative models.
  - [OpenAI Evals Framework](https://github.com/openai/evals)  
<https://github.com/openai/evals>
  - [Hugging Face Evaluation Leaderboard](https://huggingface.co/spaces/evaluate-measurement/leaderboard)  
<https://huggingface.co/spaces/evaluate-measurement/leaderboard>

## 11 Practical Projects and Applications

### 11.1 End-to-End Gen AI Projects

- Build complete applications leveraging generative AI.
  - [LangChain Practical Projects](https://github.com/gkamradt/langchain-tutorials)  
<https://github.com/gkamradt/langchain-tutorials>
  - [Microsoft TaskWeaver Framework](https://github.com/microsoft/TaskWeaver)  
<https://github.com/microsoft/TaskWeaver>

### 11.2 Industry-Specific Applications

- Learn how generative AI is transforming specific industries.
  - [McKinsey: Generative AI in Healthcare](https://www.mckinsey.com/capabilities/quantumblack/our-insights/generative-ai-in-healthcare)  
<https://www.mckinsey.com/capabilities/quantumblack/our-insights/generative-ai-in-healthcare>
  - [AWS Gen AI Case Studies](https://aws.amazon.com/solutions/case-studies/generative-ai/)  
<https://aws.amazon.com/solutions/case-studies/generative-ai/>

## 12 Ethical Considerations and Responsible AI

### 12.1 Ethical Frameworks

- Understanding ethical considerations in generative AI development.
  - [Stanford HAI Ethics Framework](https://hai.stanford.edu/sites/default/files/2020-09/AI-Ethics-Framework_HDSI.pdf)  
[https://hai.stanford.edu/sites/default/files/2020-09/AI-Ethics-Framework\\_HDSI.pdf](https://hai.stanford.edu/sites/default/files/2020-09/AI-Ethics-Framework_HDSI.pdf)
  - [UNESCO AI Ethics Guidelines](https://www.unesco.org/en/artificial-intelligence/recommendation-ethics)  
<https://www.unesco.org/en/artificial-intelligence/recommendation-ethics>

### 12.2 Bias Mitigation

- Techniques to identify and reduce bias in generative models.
  - [Ethical Considerations in NLP Research](https://huggingface.co/blog/ethics-soc-2)  
<https://huggingface.co/blog/ethics-soc-2>
  - [Anthropic's Red Teaming Approach](https://www.anthropic.com/research/red-teaming-language-models-to-reduce-harms)  
<https://www.anthropic.com/research/red-teaming-language-models-to-reduce-harms>

## 13 Future Directions in Generative AI

### 13.1 Multimodal Models

- Explore models that handle multiple modalities (text, image, audio).
  - [GPT-4V System Card](https://openai.com/research/gpt-4v-system-card)  
<https://openai.com/research/gpt-4v-system-card>
  - [Multimodal Reasoning in LLMs](https://huggingface.co/blog/multimodal-reasoning)  
<https://huggingface.co/blog/multimodal-reasoning>

## 13.2 AI Research Frontiers

- Cutting-edge research directions in generative AI.
  - [Mamba: State-Space Models](https://arxiv.org/abs/2312.11805)  
<https://arxiv.org/abs/2312.11805>
  - [Practical Guide to LLMs in Production](https://github.com/Mooler0410/LLMsPracticalGuide)  
<https://github.com/Mooler0410/LLMsPracticalGuide>

# 14 Community Resources and Continuous Learning

## 14.1 Open Source Communities

- Join communities to collaborate and learn from peers.
  - [Hugging Face Community](https://huggingface.co/)  
<https://huggingface.co/>
  - [LAION AI Organization](https://github.com/LAION-AI)  
<https://github.com/LAION-AI>

## 14.2 Conferences and Workshops

- Stay updated with the latest research through conferences.
  - [NeurIPS Conference](https://neurips.cc/)  
<https://neurips.cc/>
  - [ACL Conference Proceedings](https://aclanthology.org/venues/acl/)  
<https://aclanthology.org/venues/acl/>