

<https://techcommunity.microsoft.com/t5/ai-azure-ai-services-blog/unlocking-insights-graphrag-amp-standard-rag-in-financial/ba-p/4253311>

GraphRAG: Enhancing LLMs with Knowledge Graphs



Vijayakumar Ramdoss




Analyst | Engineer | Architect

March 2, 2025

Disclaimer: the opinions I share are solely my own and do not reflect those of my employer.

Traditional Retrieval-Augmented Generation (RAG) systems have revolutionized how we build AI applications, enabling LLMs to tap into vast amounts of external knowledge. However, these systems often struggle with complex queries and interconnected data, limiting their ability to provide comprehensive and insightful answers. GraphRAG offers a powerful solution by incorporating graph-structured knowledge, enabling more nuanced query understanding and reasoning. In this article, we'll explore how GraphRAG overcomes the limitations of traditional RAG, unlocking new possibilities for knowledge-driven AI.



Source -

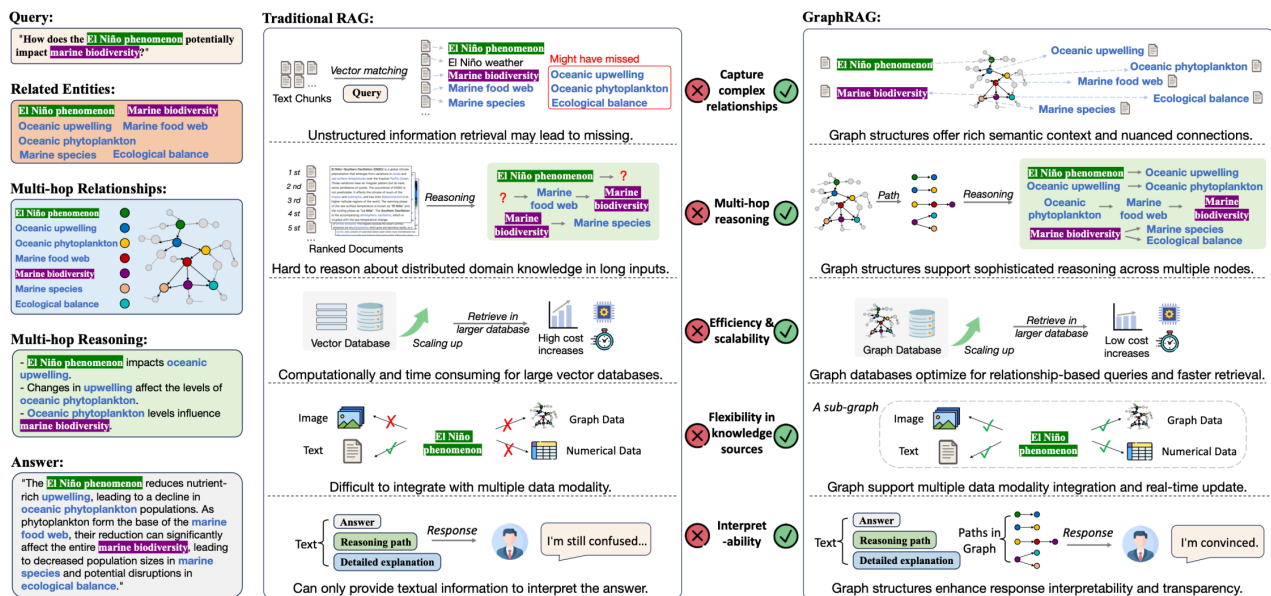
GraphRAG (Graph-based Retrieval-Augmented Generation) offers a novel method for enhancing large language models (LLMs) by integrating graph databases. This improves information retrieval and content generation. GraphRAG captures relationships and contextual information more effectively by organizing data in a graph format than traditional methods.

What is GraphRAG?

GraphRAG enhances traditional AI retrieval methods using knowledge graphs to store and retrieve structured information. Unlike traditional RAG, which relies only on vector similarity search, GraphRAG:

- Stores data as nodes (entities) and edges (relationships).
- Enables graph traversal to fetch related concepts.
- Utilizes a set of ontologies, which are rules, to help organize the information. This way, it can find hidden connections, not only the obvious ones.

The core concept of GraphRAG is that the entities in text are represented as nodes in graphs, and the relations between these entities represent the edges between the nodes. The graph is then hierarchically divided into communities and summarized into community reports.



Source -

## Why GraphRAG?

Traditional RAG systems use vector databases to retrieve relevant documents or data chunks based on embeddings and similarity search. While effective, they can struggle with:

- **Complex Relationships:** Data with intricate interconnections (e.g., knowledge graphs, hierarchical relationships).
- **Contextual Understanding:** Capturing semantic relationships between entities that span multiple dimensions.

GraphRAG addresses these challenges by leveraging graph-based data structures to enhance retrieval and reasoning capabilities.

## How GraphRAG Works

GraphRAG operates through a two-stage pipeline: Indexing and Querying.

### 1. Indexing Process

The indexing phase constructs a knowledge graph from the input corpus. This process involves several key steps:

- **Text Unit Segmentation:** The input corpus is divided into smaller, analyzable units called TextUnits.
- **Entity, Relationship, and Claims Extraction:** LLMs identify and extract all entities, relationships between them, and key claims expressed in the text from each text unit. This

extracted information is used to construct an initial knowledge graph.

- **Hierarchical Clustering:** GraphRAG uses the Leiden technique to perform hierarchical clustering on the initial knowledge graphs. Entities in each cluster are assigned to different communities for more in-depth analysis.
- **Graph Storage:** The extracted nodes and edges are stored in a graph database.

2. Querying Process

GraphRAG offers two different querying workflows tailored for different queries:

- **Global Search:** Leveraging the community summaries by reasoning about holistic questions related to the whole data corpus.
- **Local Search:** Reasoning about specific entities by fanning out to their connections.

GraphRAG's key innovation is that it structures information into a graph-based format and uses community detection to create more contextually aware responses.

GraphRAG vs. Traditional RAG

Feature	Traditional RAG	GraphRAG
Knowledge Representation	Flat, document-based, often using vector databases	<b>Graph structures</b> capturing <b>complex relationships</b>
Reasoning Capabilities	Retrieves information from chunks with anchor entities only	Enables <b>multi-hop reasoning</b> , reveals non-obvious connections
Query Understanding	Faces challenges with complex queries	Nuanced understanding, handles ambiguous queries effectively
Contextual Awareness	Limited	Captures semantic relationships across multiple dimensions
Explainability	Lacks transparency, harder to interpret	Transparent reasoning, traceable information retrieval

More info here -

Key insights from the table:

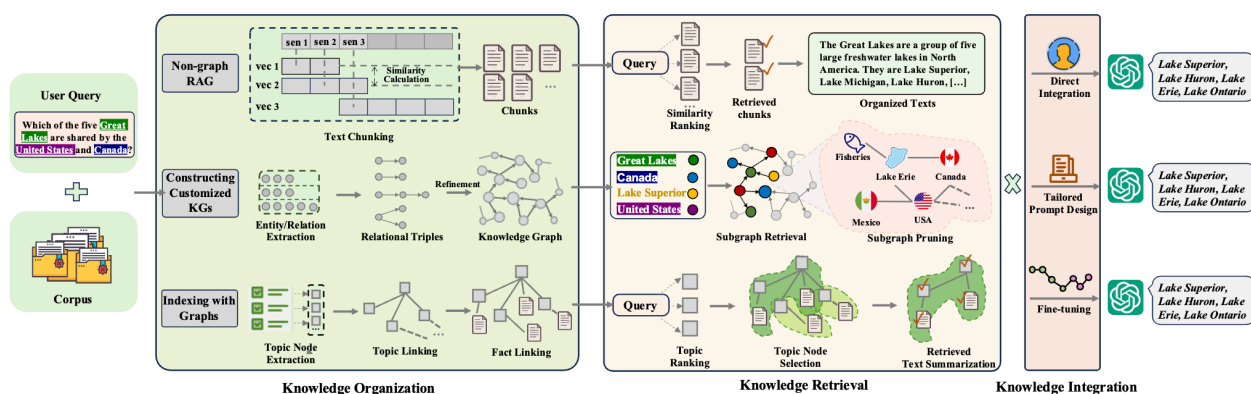
**Knowledge Representation:** GraphRAG uses graph structures to represent knowledge, capturing complex relationships between entities and concepts. Traditional RAG uses a flat, document-based representation, often relying on vector databases.

**Reasoning Capabilities:** GraphRAG enables multi-hop reasoning and can reveal non-obvious connections between different pieces of information, leading to new insights. Traditional RAG methods can only retrieve information from chunks containing anchor entities and are incapable of multi-hop reasoning.

**Query Understanding:** GraphRAG provides a more nuanced understanding of complex topics and can better handle ambiguous queries by representing multiple possible interpretations or relationships in the graph and exploring different semantic paths1.... Traditional RAG faces significant challenges in precisely answering complex queries.

**Contextual Awareness:** GraphRAG captures semantic relationships between entities across multiple dimensions. The explicit structure of knowledge graphs facilitates logic-guided chain retrieval, efficiently identifying missing facts while pruning the search space through reasoning paths1....

**Explainability:** GraphRAG provides more transparent results by tracing how information was retrieved and how relationships were used, offering more apparent, more understandable reasoning.. Traditional RAG lacks transparency and is hard to interpret.



Source - <https://arxiv.org/html/2501.13958v1>

GraphRAG improves knowledge representation compared to traditional RAG,

GraphRAG enhances knowledge representation over traditional RAG by using graph structures that capture complex relationships between entities and concepts. This approach provides a more nuanced and contextual understanding of information than traditional RAG's flat, document-based representation.

Here's a breakdown of how GraphRAG achieves this enhancement:

- **Capturing complex relationships:** GraphRAG uses graphs to model dependencies among knowledge pieces, which allows for discovering related knowledge centered around a topic. By modeling dependencies between nodes, GraphRAG enables the discovery of related knowledge pieces centered around a topic or anchor entity, ensuring comprehensive knowledge retrieval.
- **Hierarchical and multi-hop relationships:** The graph structure in GraphRAG can represent hierarchies, associations, and multi-hop relationships, offering a richer semantic context for queries and revealing non-obvious connections between different pieces of information. The explicit structure of knowledge graphs facilitates logic-guided chain retrieval, efficiently identifying missing facts while pruning the search space through reasoning paths.

- **Flexibility in knowledge sources:** GraphRAG systems can adapt to and integrate various knowledge sources, including structured databases, semi-structured data, and unstructured text. This versatility allows for a unified view of an organization's knowledge by connecting different data types and incorporating various data modalities (text, images, numerical data) into a single graph structure.
- **Context-aware reasoning:** The graph structure allows GraphRAG to consider the context of a query more effectively, leveraging semantic relationships in the graph to retrieve more relevant information. This leads to more accurate and pertinent responses than traditional RAG, which relies solely on keyword matching or vector similarity.
- **Improved handling of complex queries:** GraphRAG's structured knowledge graphs enhance the ability to answer complex queries and reduce hallucinations by providing grounding in explicitly defined relations for the answers. GraphRAG enables LLMs to generate appropriate responses accurately in context by converting unstructured knowledge bases into structured graphs to understand the information better.

### Key Benefits of GraphRAG

- **Enhanced Knowledge Representation:** Graph structures capture the relationships between entities and domain hierarchies, providing richer semantic context.
- **Improved Accuracy and Contextual Relevance:** GraphRAG better understands the context and connections between entities by leveraging structured relationships, leading to more accurate and nuanced responses.
- **Explainability:** GraphRAG allows for more transparent results by tracing how information was retrieved and how relationships were used, offering more explicit, more understandable reasoning.
- **Advanced Search Scalability:** GraphRAG efficiently navigates large and complex datasets, enabling faster and more precise searches by focusing on relevant connections between data points.

### Example Use Case for GraphRAG - Query-Focused Summarization

Consider a project focused on climate change, where our task is to address this specific question: "What impacts does climate change have on polar bears?" There's a wealth of information at our disposal, but our priority is to select the facts that specifically answer that query.

### What is Query-Focused Summarization?

Query-focused summarization means taking all the information we have and removing just the parts that answer our specific question. This helps we get straight to the point without sifting through unnecessary details.



## Scenario

Imagine we have articles, videos, and websites about climate change, and we want a summary that emphasizes its effects on polar bears. We will create a straightforward program to showcase how GraphRAG can help summarize this information.

## Code Example

We'll use Python to create a simple graph structure that includes facts about climate change and polar bears. We'll then summarize the information based on our specific question.

```
# Step 1: Create a simple Graph structure
```

```
class Node:
    def __init__(self, info):
        self.info = info
        self.connections = [] # List of nodes connected to this one

    def connect(self, other_node):
        self.connections.append(other_node)
```

```
# Step 2: Create nodes representing facts about climate change and polar bears
climate_change = Node("Climate change is causing temperatures to rise.")
polar_bears_endangered = Node("Polar bears are losing their habitat due to melting ice.")
food_sources_dwindling = Node("Ice melting reduces hunting grounds for polar bears.")
effects_on_health = Node("Climate change affects polar bear health due to temperature fluctuations.")
species_extinction = Node("If temperatures keep rising, polar bears may become extinct.")
```

```
# Step 3: Connect the nodes based on their relationships
climate_change.connect(polar_bears_endangered)
polar_bears_endangered.connect(food_sources_dwindling)
polar_bears_endangered.connect(effects_on_health)
food_sources_dwindling.connect(species_extinction)
```

```
# Step 4: Function to summarize the information based on the query
def summarize_impact_on_polar_bears(start_node):
    summary = []
    queue = [start_node] # Start exploring from the node relevant to polar bears
```

```
while queue:
    current_node = queue.pop(0)
    summary.append(current_node.info)

    for connection in current_node.connections:
        if connection not in summary: # Avoid revisiting nodes
            queue.append(connection)

return summary

# Step 5: Summarize the effects of climate change on polar bears
summary = summarize_impact_on_polar_bears(climate_change)

# Step 6: Print the summary
print("Summary of the Effects of Climate Change on Polar Bears:")
for fact in summary:
    print("- " + fact)
```

### Explanation of the Code

1. Graph Structure: We create a Node class to represent pieces of information (facts).
2. Create Facts: We make nodes for different facts about climate change and polar bears. Each node is like a piece of information connected to others.
3. Connect Nodes: We connect these nodes to show their relationships, which helps us understand how they relate to the main topic.
4. Summarize Function: The `summarize_impact_on_polar_bears` function pulls together all the facts about how climate change affects polar bears.
5. Output the Summary: Finally, we print out the summary, which gives us the relevant information without extra noise.

When we execute the code, it will produce a summary of how climate change affects polar bears:

```
Summary of the Effects of Climate Change on Polar Bears:
- Climate change is causing temperatures to rise.
- Polar bears are losing their habitat due to melting ice.
- Ice melting reduces hunting grounds for polar bears.
```



- Climate change affects polar bear health due to temperature fluctuations.
- If temperatures keep rising, polar bears may become extinct.

This code simulates how GraphRAG gathers relevant information for our question, effectively organizing it to answer our inquiry about polar bears and climate change.

### Practical Applications of GraphRAG

- **Pharmaceutical Research:** Retrieve and explain drug mechanisms, interactions, and side effects.
- **Knowledge Management:** Generate insights from corporate knowledge graphs.
- **Recommendation Systems:** Recommend products or content based on user behavior graphs.
- **Legal and Compliance:** Extract and summarize regulations, case laws, and their interconnections.
- **Financial Services:** Transformative tool in customer relationship management and risk assessment.
- **Healthcare:** Resolve intricate queries regarding treatment protocols by accessing interconnected data sources efficiently.

### Implementation Challenges of GraphRAG

- **Graph Construction:** Creating a structured knowledge graph from unstructured data requires attention to detail and rigorous processes.
- **Scalability:** Knowledge graphs grow in complexity with the size of the knowledge base, potentially causing scalability issues. Optimizing the amount of related information to retrieve can lead to overload during information retrieval.
- **Maintenance Overhead:** Knowledge graphs need constant updating with new information and changes in existing data.
- **Query Generation:** Generating effective queries for querying these graphs is a primary challenge.
- **Reasoning Boundary:** Deciding how much information to retrieve based on queries is an explicit limitation.

### Future Research Directions for GraphRAG

- Improved Computational Efficiency and Scalability: Future LLMs are expected to have broader and deeper knowledge across multiple domains.
- Hybrid Approaches: RAG applications will likely adopt hybrid approaches, combining extractive queries handled by text2emb RAG with global abstractive queries using knowledge graphs.
- Automated Graph Construction and Maintenance: Automated graph construction and maintenance methods are needed to reduce manual effort and ensure up-to-date knowledge graphs.
- Explainable AI: Future research should focus on enhancing the interpretability of results through graph structures.

By integrating advanced graph theory with AI-driven processes, GraphRAG delivers exceptional accuracy and contextual depth. As organizations adapt to an increasingly connected data landscape, the adoption of GraphRAG presents a crucial opportunity to harness the power of interconnected insights that drive informed decisions.

## Conclusion

GraphRAG enhances data retrieval and generation by integrating graph theory with AI, enhancing accuracy and context. It addresses issues of traditional Retrieval-Augmented Generation (RAG) systems with complex queries and large datasets. GraphRAG improves precision, context-awareness, and explainability by utilizing structured relationships in knowledge graphs. Advances in graph construction and data quality are vital for enhancing retrieval systems and maximizing GraphRAG's potential in data workflows.

## Reference:

<https://github.com/DEEP-PolyU/Awesome-GraphRAG>

<https://arxiv.org/abs/2501.13958>

<https://arxiv.org/pdf/2404.16130>